



Amazon Web Services Data Engineering Immersion Day

Lab 1. Hydrating the Data Lake with Glue Streaming ETL
March 2021

Table of Contents

Introduction	2
Go To Lab Environment	3
Setup Streaming Data Generator.....	5
Create Kinesis Data Stream	5
Create Table for Kinesis Stream Source in Glue Data Catalog	6
Create and trigger Glue Stream job.....	9
Trigger stream data from KDG.....	12
Verify the Glue stream job	14
Create Glue Crawler for the transformed data.....	15
Trigger abnormal transaction data from KDG	19
Detect Abnormal Transactions using Ad-Hoc query from Athena	21

Introduction

This lab will guide you to understand AWS Glue Streaming ETL feature. You will start with hydrating your Data Lake from Amazon Kinesis Data Generator (KDG). The final outcome is to query the Data Lake in near real-time.



In this lab you will complete the following tasks:

1. Setup a Streaming Data Generator for Kinesis
2. Create Glue Streaming job
3. Query the data stream in Athena

If you'd like to run the workshop on your own after the AWS hosted event, please follow the lab instruction here: <https://github.com/aws-samples/data-engineering-for-aws-immersion-day>

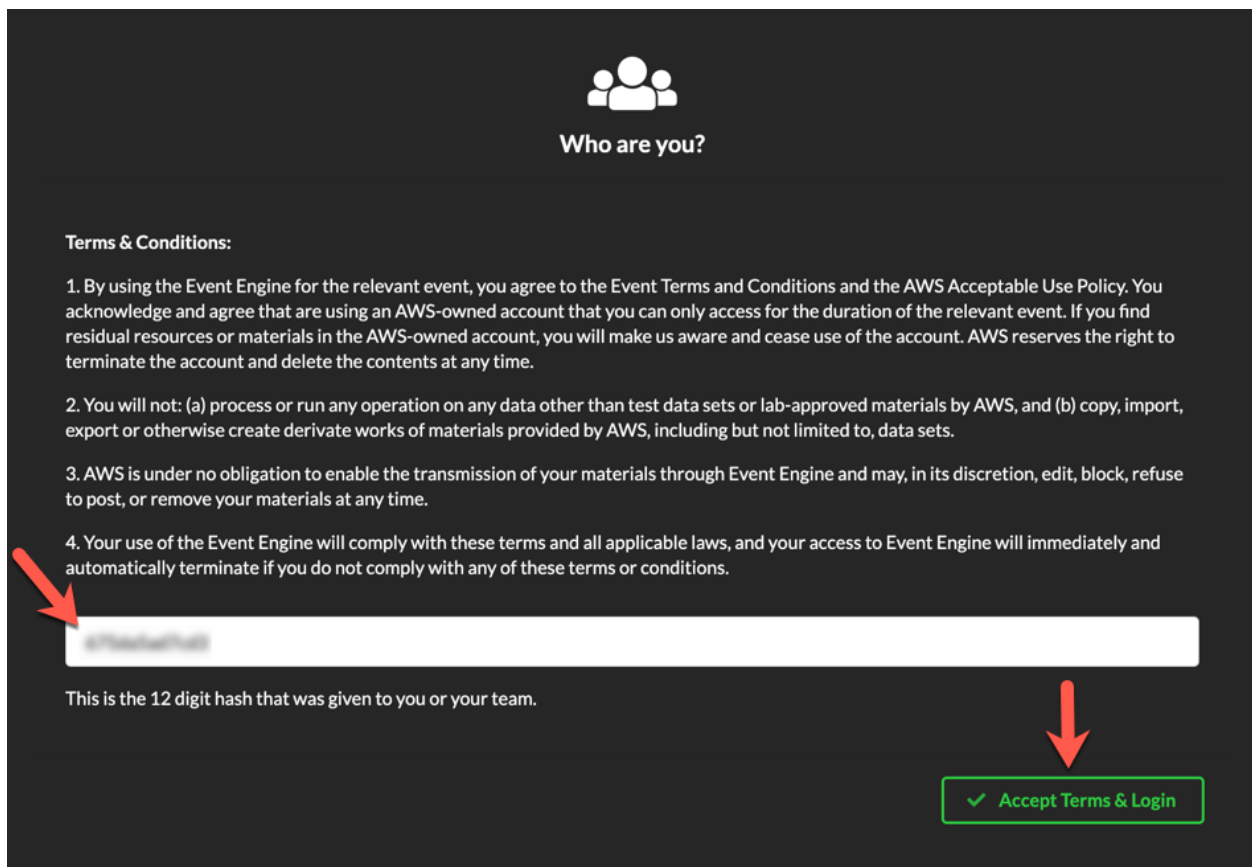
Go To Lab Environment

Please skip this section if you are running the lab on your own AWS account.

Today, you are attending a formal event and you will have been sent your access details beforehand. If in the future you might want to perform these labs in your own AWS environment by yourself, you can follow instructions on GitHub - <https://github.com/aws-samples/data-engineering-for-aws-immersion-day>.

A 12-character access code (or 'hash') is the access code that grants you permission to use a dedicated AWS account for the purposes of this workshop.

1. Go to <https://dashboard.eventengine.run/>, enter the access code and click Proceed:



Who are you?

Terms & Conditions:

1. By using the Event Engine for the relevant event, you agree to the Event Terms and Conditions and the AWS Acceptable Use Policy. You acknowledge and agree that are using an AWS-owned account that you can only access for the duration of the relevant event. If you find residual resources or materials in the AWS-owned account, you will make us aware and cease use of the account. AWS reserves the right to terminate the account and delete the contents at any time.
2. You will not: (a) process or run any operation on any data other than test data sets or lab-approved materials by AWS, and (b) copy, import, export or otherwise create derivate works of materials provided by AWS, including but not limited to, data sets.
3. AWS is under no obligation to enable the transmission of your materials through Event Engine and may, in its discretion, edit, block, refuse to post, or remove your materials at any time.
4. Your use of the Event Engine will comply with these terms and all applicable laws, and your access to Event Engine will immediately and automatically terminate if you do not comply with any of these terms or conditions.

This is the 12 digit hash that was given to you or your team.


✓ Accept Terms & Login

2. On the Team Dashboard web page, you will see a set of parameters that you will need during the labs. Best to save them to a text file locally, alternatively you can always go to this page to review them. Replace the parameters with the corresponding values from here where indicated in subsequent labs:

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Because you're at a formal event, some AWS resources have been pre-deployed for your convenience, for example:

- S3 Bucket, IAM roles etc

 Modules

Environment Setup

[Readme](#)

Outputs:

S3 Bucket name
mod-3fccddd609114925-dmslabs3bucket-1ngcgzzcnd15u

BusinessAnalystUser
mod-3fccddd609114925-BusinessAnalystUser-MBOXFZLQLOXX

DMSLabRoleS3 ARN
arn:aws:iam::377243295828:role/mod-3fccddd609114925-DMSLabRoleS3-O2VT1RSN43SG

Glue Lab Role
mod-3fccddd609114925-GlueLabRole-YLTJA13WW6WT


S3BucketWorkgroupA
mod-3fccddd609114925-s3bucketworkgroupa-tbon3m1mkunh

S3BucketWorkgroupB
mod-3fccddd609114925-s3bucketworkgroupb-18ygl8nfp8ead

WorkgroupManagerUser
mod-3fccddd609114925-WorkgroupManagerUser-5IVE0UQNIBG4

3. On the Team Dashboard, please click AWS Console to log into the AWS Management Console:

Team Dashboard

 Event

[AWS Console](#) [SSH Key](#)

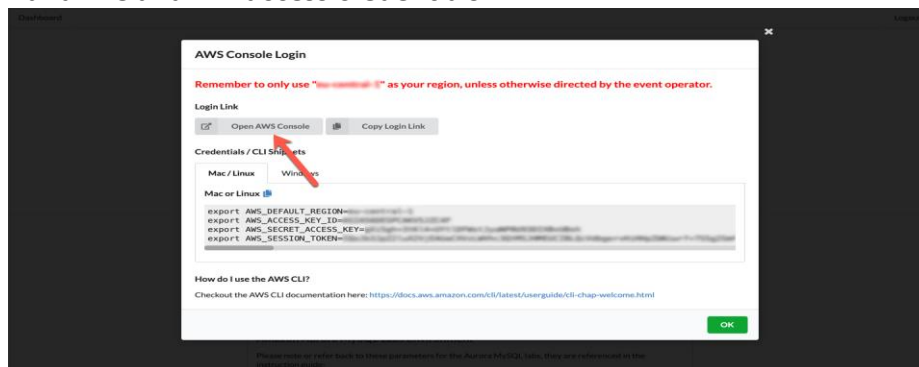
Event Data Engineering Immersion Day - Test

Team Name:

Event ID: d2302d4ae9ff4ea2857846b74f7de7e2

Team ID: 1c2f7ad7ec044b0b8276f917c5983133

4. Click Open AWS Console. For the purposes of this workshop, you will not need to use command line and API access credentials:



Once you have completed these steps, you can continue with the rest of this lab.

Setup Streaming Data Generator

We need an Amazon Kinesis Data Generator (Amazon KDG) to simulate the streaming data. If you have set up Amazon KDG with Kinesis Clickstream Lab, you should be able to reuse the tool. Otherwise, please follow the instruction in [Streaming Data Prelab](#) to launch the CloudFormation template, in order to set up your Amazon Kinesis Data Generator.

After the KDG setup is completed, you can find a URL from the output tab of the Streaming Data Prelab CloudFormation Stack, with the key name KinesisDataGeneratorUrl. Make sure you can login to the console using the username and password you provided when launching your prelab CloudFormation template. Bookmark the URL for further use.

Create Kinesis Data Stream

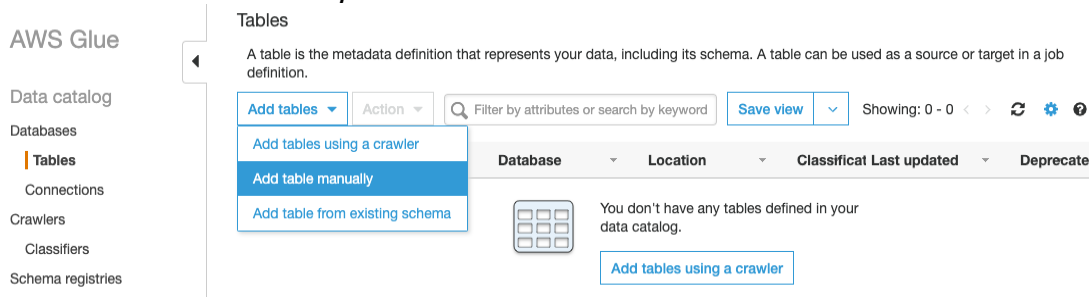
1. Navigate to [AWS Kinesis console](#) by using this link
2. Click **"Create data stream"**
3. Put **"TicketTransactionStreamingData"** as data stream name and put number of open shards as 2, then click **"Create data stream"**.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

The screenshot shows the 'Create data stream' page in the Amazon Kinesis console. The breadcrumb navigation is 'Amazon Kinesis > Data streams > Create data stream'. The main heading is 'Create a data stream' with an 'Info' link. Below this is the 'Data stream configuration' section, which includes a 'Data stream name' field containing 'TicketTransactionStreamingData'. A note states: 'Acceptable characters are uppercase and lowercase letters, numbers, underscores, hyphens and periods.' The next section is 'Data stream capacity' with an 'Info' link and a 'Request limit increase' button. It explains that data records are stored in Kinesis Data Stream and a shard is a uniquely identified sequence of data records. Under 'Shard estimator', it shows 'Number of open shards' set to 2, with a note that each shard ingests up to 1 MiB/second and 1000 records/second and emits up to 2 MiB/second. It also shows 'Total data stream capacity' calculated based on the number of shards. The 'Write' capacity is '2 MiB/second, 2000 Data records/second' and the 'Read' capacity is '4 MiB/second'. At the bottom, there are 'Cancel' and 'Create data stream' buttons.

Create Table for Kinesis Stream Source in Glue Data Catalog

1. Navigate to [AWS Glue](#) console by using this link
2. On the AWS Glue menu, select Tables



3. Put **TicketTransactionStreamData** as the table name
4. Click Add database and put **tickettransactiondatabase** as the database name, and click create.

Add database

The screenshot shows the 'Add database' form in the AWS Glue console. It has a 'Database name' field containing 'tickettransactiondatabase'. Below this is a section for 'Description and location (optional)' which is currently empty.

5. Using drop down to select the database we just created, and click Next

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Set up your table's properties

Table name
TicketTransactionStreamData

Database ⓘ
tickettransactiondatabase

Add database

Description (optional)

Next

6. Select **Kinesis** as the source, put stream name as **TicketTransactionStreamingData**, for the kinesis source URL, put in <https://kinesis.us-east-1.amazonaws.com>, click **Next**.

Add a data store

Select the type of source

☐ S3
☒ Kinesis
☐ Kafka

Stream name
TicketTransactionStreamingData

Kinesis source URL
https://kinesis.us-east-1.amazonaws.com

Back Next

7. Choose **JSON** as the incoming data format, as we will trigger JSON payload from Kinesis Data Generator in following steps. Click **Next**.

Choose a data format

Classification

☐ CSV
☒ JSON
☐ ORC
☐ Parquet
☐ Avro
☐ Grok

Choose the format of the data in your table.

Back Next

8. Leave the schema as empty, as we will enable schema detection feature when defining glue stream job. Click **Next**.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Add table

Table properties

Name:
TicketTransactionStreamData
Database:
tickettransactiondatabase

Data store

JSON

Data format

JSON

Schema

Review

Define a schema

Add column

Showing: 0 - 0 of 0 < >

Column name	Data type	Key	Comment
You don't have any columns defined yet. To create a column, choose Add column			

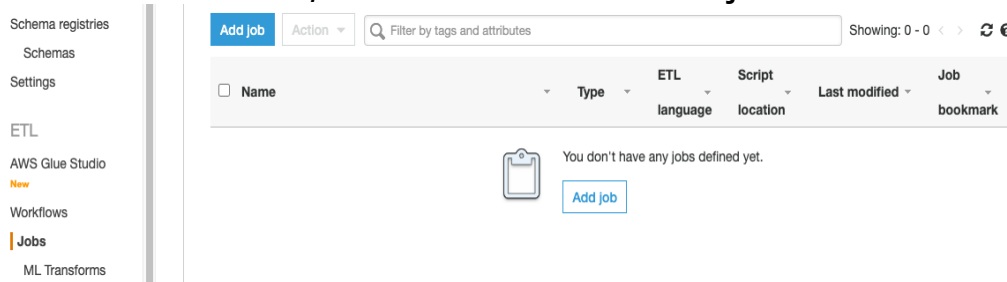
Back

Next

- Review all the details and click **Finish**.

Create and trigger Glue Stream job

1. Navigate to [AWS Glue console](#)
2. On the AWS Glue menu, select **Jobs** and then click **Add job**



3. Put **TicketTransactionStreamingJob** as the job name, select the IAM role with "GlueLabRole" in the name. For job type, use dropdown list, select **Spark Streaming**;

Configure the job properties

Name

IAM role ⓘ

 ⓘ

Ensure that this role has permission to your Amazon S3 sources, targets, temporary directory, scripts, and any libraries used by the job. [Create IAM role.](#)

Type

☒ A proposed script generated by AWS Glue ⓘ
 ☐ An existing script that you provide
 ☐ A new script to be authored by you

Script file name

S3 path where the script is stored

 ⓘ

Temporary directory ⓘ

 ⓘ

4. leave the rest configurations as is and click **Next**.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

5. For Data source, select the data source **tickettransactionstreamdata**, then click **Next**.

Choose a data source

Filter by attributes or search by keyword

Showing: 1 - 1

Name	Database	Location	Classification
tickettransactionstreamdata	tickettransactiondatabase	TicketTransactionStreaming...	json

6. In Data target, select **Create tables in your data target**. In Data store dropdown list, select **Amazon S3**. Select **Parquet** format from dropdown list.

Choose a data target

☒ Create tables in your data target
☐ Use tables in the data catalog and update your data target

Data store
Amazon S3

Format
Parquet

Connection
- Select one -

Add connection

Target path
s3://bucket/prefix/object

7. Click the **folder** button next to **Target path** to select a S3 bucket. From the pop-up window, select a S3 bucket name that contains "**dmslabs3bucket**".

Choose S3 path

S3

- aws-glue-scripts-...-us-east-1
- aws-glue-temporary-488748888187-us-east-1
- mod-...-dmslabs3bucket-1...
- ...
- ...

8. Make sure you add a path at the end **/TicketTransactionStream**

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Choose a data target

☒ Create tables in your data target
☐ Use tables in the data catalog and update your data target

Data store
Amazon S3

Format
Parquet

Connection
- Select one -
[Add connection](#)

Target path
dmslabs3bucket-...s/TicketTransactionStream

[Back](#) [Next](#)

9. Make sure you select **Automatically detect schema of each record**, then click **Save job and edit script**.

Output Schema Definition

☒ Automatically detect schema of each record
The output schema will be inferred from the input stream.

☐ Specify output schema for all records
Use an Apply Mapping transform to define the output schema.

[Back](#) [Save job and edit script](#)


10. Review the generated script, click **Save** and then quit the editor.

Job: TicketTransactionStreamingJob [Action](#) [Save](#) [Run job](#)

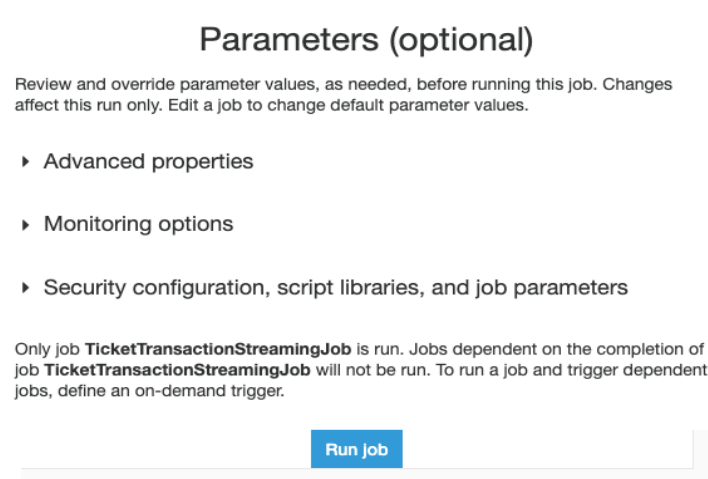
```
1 import sys
2 from aws glue.transforms import *
3 from aws glue.utils import getResolvedOptions
4 from pyspark.context import SparkContext
5 from aws glue.context import GlueContext
6 from aws glue.job import Job
7 from pyspark.sql import DataFrame, Row
8 import datetime
9 from aws glue import DynamicFrame
10
11 ## @params: [JOB_NAME]
12 args = getResolvedOptions(sys.argv, ['JOB_NAME'])
13
14 sc = SparkContext()
15 glueContext = GlueContext(sc)
16 spark = glueContext.spark_session
17 job = Job(glueContext)
18 job.init(args['JOB_NAME'], args)
19 ## @type: DataSource
```

11. Select the **TicketTransactionStreamingJob** we just created, from the Action dropdown list, select **Run job**.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL



12. Leave the optional parameters as default and click **Run job** to trigger the Glue Stream Job



Trigger stream data from KDG

1. Launch KDG using the URL you bookmarked from the lab setup, login using the username and password you specified when deploying the CloudFormation stack.
2. Make sure you select **us-east-1** region, from the dropdown list, select the **TicketTransactionStreamingData** as the target Kinesis stream, leave Records per second as default (**100** records per second); for the record template, type in **NormalTransaction** as the payload name, and copy the template payload as below:

```
{
  "customerId": "{{random.number(50)}}",
```

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

```
"transactionAmount": {{random.number(
  {
    "min":10,
    "max":150
  }
)}}},
"sourceIp": "{{internet.ip}}",
"status": "{{random.weightedArrayElement(
  "weights": [0.8,0.1,0.1],
  "data": ["OK","FAIL","PENDING"]
  )}}",
"transactionTime": "{{date.now}}"
}
```

Region: us-east-1

Stream/delivery stream: TicketTransactionStreamingData

Records per second: Constant **Periodic**

100

Compress Records ☐

Record template ⓘ

NormalTransaction **Template 2** Template 3 Template 4 Template 5

NormalTransaction

```
{
  "customerId": "{{random.number(50)}}",
  "transactionAmount": {{random.number(
    {
      "min":10,
      "max":150
    }
)}}},
  "sourceIp": "{{internet.ip}}",
  "status": "{{random.weightedArrayElement(
    "weights": [0.8,0.1,0.1],
    "data": ["OK","FAIL","PENDING"]
  )}}",
  "transactionTime": "{{date.now}}"
}
```


Send data Test template

To learn more about what the payload will look like when sending from KDG simulator, refer to the document as this link,

<https://awslabs.github.io/amazon-kinesis-data-generator/web/help.html>


Lab 1. Hydrating the Data Lake with Glue Streaming ETL

3. Click **Send data** to trigger the simulated ticket purchasing transaction streaming data.



Verify the Glue stream job

1. Navigate to **Amazon S3** console by using this link <https://s3.console.aws.amazon.com/s3/home?region=us-east-1>
2. Navigate to the S3 bucket path we've set as Glue Stream Job target, note the folder structure of the processed data.



3. Check the folder content using current date and time as the folder name. Verify that the streaming data has been transformed into parquet format and persisted into corresponding folders.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Amazon S3 > mod-[redacted]-dmslabsbucket-[redacted] > TicketTransactionStreamingData/ > ingest_year=2021/ > ingest_month=02/ > ingest_day=01/ > ingest_hour=10/

ingest_hour=10/ Copy S3 URI

Objects | Folder properties

Objects (8)
Objects are the fundamental entities stored in Amazon S3. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Refresh Delete Actions Create folder Upload

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	part-00000-1442bf92-c767-4277-bda8-8d54578fab3d-c000.snappy.parquet	parquet	February 1, 2021, 18:31:45 (UTC+08:00)	25.2 KB	Standard
<input type="checkbox"/>	part-00000-b97c00da-9a5a-4d81-9613-16e37e75c852-c000.snappy.parquet	parquet	February 1, 2021, 18:35:06 (UTC+08:00)	64.3 KB	Standard
<input type="checkbox"/>	part-00000-e809524f-0576-4b4a-a213-cff4a787d3ba-c000.snappy.parquet	parquet	February 1, 2021, 18:30:14 (UTC+08:00)	3.4 KB	Standard
<input type="checkbox"/>	part-00000-fd8312ce-9484-42be-96db-82ef48f775d9-c000.snappy.parquet	parquet	February 1, 2021, 18:33:26 (UTC+08:00)	64.2 KB	Standard
<input type="checkbox"/>	part-00001-1442bf92-c767-4277-bda8-8d54578fab3d-c000.snappy.parquet	parquet	February 1, 2021, 18:31:45 (UTC+08:00)	24.9 KB	Standard
<input type="checkbox"/>	part-00001-b97c00da-9a5a-4d81-9613-16e37e75c852-c000.snappy.parquet	parquet	February 1, 2021, 18:35:07 (UTC+08:00)	63.5 KB	Standard
<input type="checkbox"/>	part-00001-e809524f-0576-4b4a-a213-cff4a787d3ba-c000.snappy.parquet	parquet	February 1, 2021, 18:30:14 (UTC+08:00)	3.4 KB	Standard
<input type="checkbox"/>	part-00001-fd8312ce-9484-42be-96db-82ef48f775d9-c000.snappy.parquet	parquet	February 1, 2021, 18:33:26 (UTC+08:00)	62.5 KB	Standard

Create Glue Crawler for the transformed data


1. Navigate to [AWS Glue console](#)
2. On the AWS Glue menu, select **Crawlers** and click **Add crawler**.

AWS Glue

- Data catalog
- Databases
- Tables
- Connections
- Crawlers**
- Classifiers
- Schema registries

Crawlers A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then exports the schema to the Data Catalog.

Add crawler Run crawler Action

<input type="checkbox"/>	Name	Schedule	Status	Logs	Last runtime
 You don't have any crawlers yet. Add crawler					

3. Put **TicketTransactionParquetDataCrawler** as the name of the crawler, click **Next**.

Add information about your crawler

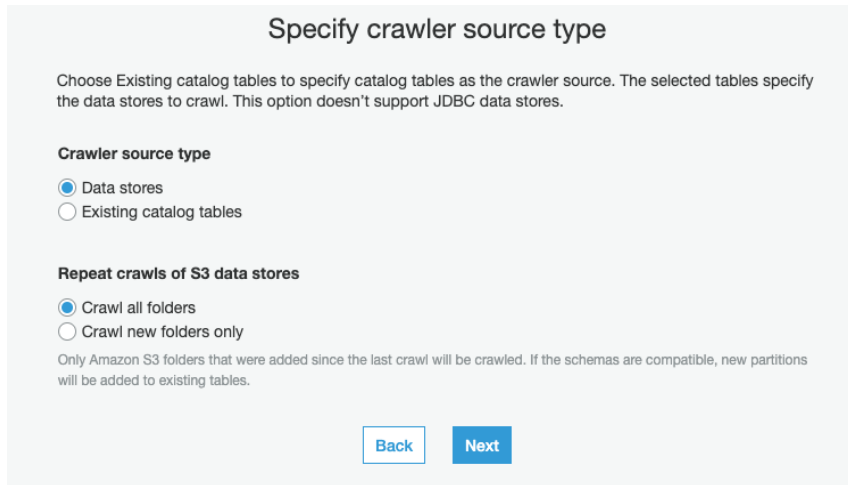
Crawler name

Tags, description, security configuration, and classifiers (optional)

Next

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

4. Leave the default to specify **Data stores** as Crawler source type and **Crawl all folders**, click **Next**.



Specify crawler source type

Choose Existing catalog tables to specify catalog tables as the crawler source. The selected tables specify the data stores to crawl. This option doesn't support JDBC data stores.

Crawler source type

☒ Data stores
☐ Existing catalog tables

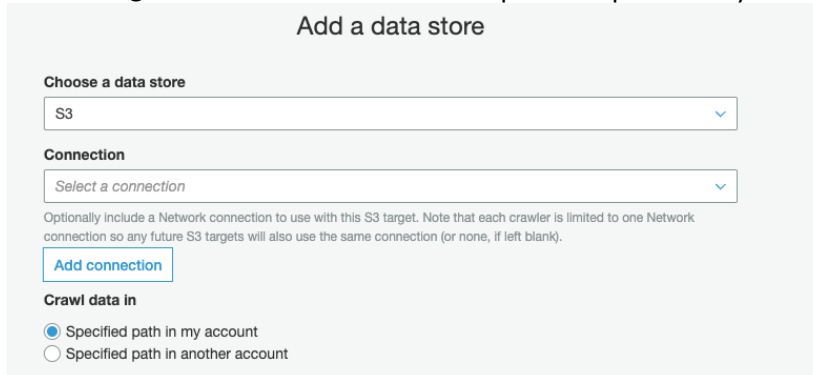
Repeat crawls of S3 data stores

☒ Crawl all folders
☐ Crawl new folders only

Only Amazon S3 folders that were added since the last crawl will be crawled. If the schemas are compatible, new partitions will be added to existing tables.

[Back](#) [Next](#)

5. Choose S3 as data store and choose Specified path in my account.



Add a data store

Choose a data store

S3

Connection

Select a connection

Optionally include a Network connection to use with this S3 target. Note that each crawler is limited to one Network connection so any future S3 targets will also use the same connection (or none, if left blank).

[Add connection](#)

Crawl data in

☒ Specified path in my account
☐ Specified path in another account

6. Click the icon next to Include path input to select the S3 bucket. Make sure you select the folder **TicketTransactionStreamingData**. Click **Select**.



Choose S3 path

S3

- ☐ mod-...
- ☐ mod-...
- ☐ mod-...
- ☐ mod-...
- ☐ mod-...
- ☐ mod-...
- ☒ TicketTransactionStreamingData
- ☐ mod-...
- ☐ mod-...

[Select](#)

7. Expand the **Exclude patterns**, put **checkpoint/**** to exclude the data in checkpoint folder. Review the current input and click **Next**.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Add a data store

Choose a data store

S3

Connection

Select a connection

Optionally include a Network connection to use with this S3 target. Note that each crawler is limited to one Network connection so any future S3 targets will also use the same connection (or none, if left blank).

Add connection

Crawl data in

☒ Specified path

Include path

s3://mod-*XXXXXXXXXX*-dmslabs3bucket-*XXXXXXXXXX*/TicketTransactionStreamingC

All folders and files contained in the include path are crawled. For example, type s3://MyBucket/MyFolder/ to crawl all objects in MyFolder within MyBucket.

▼ Exclude patterns (optional)

Exclude patterns

checkpoint/**

glob pattern

The exclude pattern is relative to the include path. Objects that match the exclude pattern are not crawled. For example, with include path s3://mybucket/ and exclude pattern, mydir/**, then all objects in the include path below the mydir directory are skipped. In this example, any object whose path matches s3://mybucket/mydir/** is not crawled. For more information about patterns, see [Cataloging Tables with a Crawler.](#)

Back Next

8. Select No to indicate no other data store needed, then click Next.

Add another data store

☐ Yes

☒ No

Back Next

9. Choose an existing IAM role, using the dropdown list to select the role with **GlueLabRole** in the name, click **Next**.

Choose an IAM role

The IAM role allows the crawler to run and access your Amazon S3 data stores. [Learn more](#)

☐ Update a policy in an IAM role

☒ Choose an existing IAM role

☐ Create an IAM role

IAM role

mod-*XXXXXXXXXX*-GlueLabRole-1*XXXXXXXXXX*

This role must provide permissions similar to the AWS managed policy, **AWSGlueServiceRole**, plus access to your data stores.

- s3://mod-*XXXXXXXXXX*-dmslabs3bucket-*XXXXXXXXXX*/TicketTransactionStreamingData/ingest_year=2021

You can also create an IAM role on the [IAM console](#).

Back Next

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

- As the data is partitioned to hour, so we set the crawler to run every hour to make sure newly added partition is added. Click **Next**.

Create a schedule for this crawler

Frequency

Hourly

Start Minute

00

Back Next

- Using the dropdown list to select **tickettransactiondatabase** as the output database, use **parquet_** as the prefix for the table, click **Next**.

Configure the crawler's output

Database

tickettransactiondatabase

Add database

Prefix added to tables (optional)

parquet_

Grouping behavior for S3 data (optional)

Configuration options (optional)

Back Next

- Review the crawler configuration and click **Finish** to create the crawler.
- Once the crawler is created, select the crawler and click **Run crawler** to trigger the first run.

[user preferences](#)

Add crawler Run crawler Action Filter by tags and attributes Showing: 1 - 1

<input checked="" type="checkbox"/>	Name	Schedule	Status	Logs	Last runtime	Median runtime	Tables updated	Tables added
<input checked="" type="checkbox"/>	TicketTransactionParquetDataCrawler	At 00 minutes past ...	Ready		0 secs	0 secs	0	0

- When crawler job stopped, go to Glue Data catalog, under Tables, verify that **parquet_tickettransactionstreamingdata** table is listed.

<input type="checkbox"/>	Name	Database	Location	Classification
<input type="checkbox"/>	parquet_tickettransactionstreamingdata	tickettransactiondatabase	s3://mod-3fccddd609114925-d...	parquet
<input type="checkbox"/>	tickettransactionstreamdata	tickettransactiondatabase	TicketTransactionStreamingData	json

- Click the **parquet_tickettransactionstreamingdata** table, verify that Glue has correctly identified the streaming data format while transforming source data from Json format to Parquet.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Tables	>	parquet_tickettransactionstreamingdata	Last updated 1 Feb 2021 07:39 PM	Table	Version (Current version)	▼
Edit table	Delete table		View partitions	Compare versions	Edit schema	
Name	parquet_tickettransactionstreamingdata					
Description						
Database	tickettransactiondatabase					
Classification	parquet					
Location	s3://mod-5f6e6e6e6e6e-dmslabs3bucket-[REDACTED]/TicketTransactionStreamingData/					
Connection						
Deprecated	No					
Last updated	Mon Feb 01 19:39:02 GMT+800 2021					
Input format	org.apache.hadoop.hive.q.l.o.parquet.MapredParquetInputFormat					
Output format	org.apache.hadoop.hive.q.l.o.parquet.MapredParquetOutputFormat					
Serde serialization lib	org.apache.hadoop.hive.q.l.o.parquet.serde.ParquetHiveSerDe					
Serde parameters	serialization.format 1 sizeKey 1121888 objectCount 28 UPDATED_BY_CRAWLER TicketTransactionParquetDataCrawler CrawlerSchemaSerializerVersion 1.0 recordCount 85300 averageRecordSize 20 exclusions ["s3://mod-5f6e6e6e6e6e-dmslabs3bucket-[REDACTED]/TicketTransactionStreamingData/checkpoint/*"] CrawlerSchemaDeserializerVersion 1.0 compressionType none typeOfData file					
Table properties						

Trigger abnormal transaction data from KDG

1. Keep the KDG streaming data running, open another browser and launch KDG using the URL you bookmarked from the lab setup, login using the username and password you provided when launching the CloudFormation template.
2. Make sure you select **us-east-1** region, from the dropdown list, select the **TicketTransactionStreamingData** as the target Kinesis stream, put Records per second as 1; click Template 2, and prepare to copy abnormal transaction data,

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Region

Stream/delivery stream

Records per second

☐ Constant ☒ Periodic

Compress Records ☐

Record template ☒ NormalTransaction ☐ Template 2 ☐ Template 3 ☐ Template 4 ☐ Template 5

- for the record template, type in **NormalTransaction** as the payload name, and copy the template payload as below:

```
{
  "customerId": "{{random.number(50)}}",
  "transactionAmount": {{random.number(
    {
      "min":10,
      "max":150
    }
  )}},
  "sourceIp": "221.233.116.256",
  "status": "{{random.weightedArrayElement({
    "weights": [0.8,0.1,0.1],
    "data": ["OK","FAIL","PENDING"]
  })}}",
  "transactionTime": "{{date.now}}"
}
```

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Region: us-east-1

Stream/delivery stream: TicketTransactionStreamingData

Records per second: Constant Periodic

Compress Records ☐

Record template NormalTransaction AbnormalTransaction Template 3 Template 4 Template 5

AbnormalTransaction

```
{
  "customerId": "{{random.number(50)}}",
  "transactionAmount": {{random.number(
    {
      "min":10,
      "max":150
    }
  )}},
  "sourceIp" : "221.233.116.256",
  "status": "{{random.weightedArrayElement({
    "weights" : [0.8,0.1,0.1],
    "data": ["OK","FAIL","PENDING"]
  })}}",
  "transactionTime": "{{date.now}}"
}
```

Send data Test template

4. Click Send data to simulate abnormal transactions (1 transaction per second all from the same source IP address).

Detect Abnormal Transactions using Ad-Hoc query from Athena

1. Navigate to **AWS Athena** console by using this link <https://console.aws.amazon.com/athena/home?region=us-east-1>
2. Make sure you select **AwsDataCatalog** as Data source and **tickettransactiondatabase** as the database, refresh to make sure the **parquet_tickettransactionstreamingdata** is showing in the table list.

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

Data source

Connect data source

AwsDataCatalog

Database

tickettransactiondatabase

Filter tables and views...

▼ Tables (1)

Create table

▶ parquet_tickettransactionstreamingdata (Partitioned)

▼ Views (0)

Create view

You have not created any views. To create a view, run a query and click "Create view from query"

3. Copy query as below, this is to query last hour the number of transactions by sourceip. You should see there's large number of transactions from the same sourceip.

```
SELECT count(*) as numberOfTransactions, sourceip
FROM "tickettransactiondatabase"."parquet_tickettransactionstreamingdata"
WHERE ingest_year='2021'
AND cast(ingest_year as bigint)=year(now())
AND cast(ingest_month as bigint)=month(now())
AND cast(ingest_day as bigint)=day_of_month(now())
AND cast(ingest_hour as bigint)=hour(now())
GROUP BY sourceip
Order by numberOfTransactions DESC;
```

```
1 SELECT count(*) as numberOfTransactions, sourceip
2 FROM "tickettransactiondatabase"."parquet_tickettransactionstreamingdata"
3 WHERE ingest_year='2021'
4 AND cast(ingest_year as bigint)=year(now())
5 AND cast(ingest_month as bigint)=month(now())
6 AND cast(ingest_day as bigint)=day_of_month(now())
7 AND cast(ingest_hour as bigint)=hour(now())
8 GROUP BY sourceip
9 Order by numberOfTransactions DESC;
10
```

Run query

Save as

Create

(Run time: 2.47 seconds, Data scanned: 1.14 MB)

Format query

Clear

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Athena engine version 1 [Release versions](#)

Results

	numberOfTransactions ▼	sourceip ▼
1	4468	221.233.116.256
2	2	192.45.173.73
3	2	120.233.79.63
4	1	2.237.235.165
5	1	166.88.59.49
6	1	144.70.141.118
7	1	123.45.98.210
8	1	14.46.63.97

Lab 1. Hydrating the Data Lake with Glue Streaming ETL

- Copy query as below, this is to further check if the transaction details from the same source IP. The query verified that the request is coming from same IP but with different customer id, so it's verified as abnormal transactions.

```
SELECT *
FROM "tickettransactiondatabase"."parquet_tickettransactionstreamingdata"
WHERE ingest_year='2021'
AND cast(ingest_year as bigint)=year(now())
AND cast(ingest_month as bigint)=month(now())
AND cast(ingest_day as bigint)=day_of_month(now())
AND cast(ingest_hour as bigint)=hour(now())
AND sourceip='221.233.116.256'
limit 100;
```

```
1 SELECT *
2 FROM "tickettransactiondatabase"."parquet_tickettransactionstreamingdata"
3 WHERE ingest_year='2021'
4 AND cast(ingest_year as bigint)=year(now())
5 AND cast(ingest_month as bigint)=month(now())
6 AND cast(ingest_day as bigint)=day_of_month(now())
7 AND cast(ingest_hour as bigint)=hour(now())
8 AND sourceip='221.233.116.256'
9 limit 100;
```

Run query

Save as

Create

(Run time: 3.49 seconds, Data scanned: 650.97 KB)

Format query

Clear

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Athena engine version 1

[Release versions](#)

Results

	customerid	sourceip	status	transactionamount	transactiontime	ingest_year	ingest_month	ingest_day	ingest_hour
1	4	221.233.116.256	OK	117	2021-02-01T20:31:46+08:00	2021	02	01	12
2	26	221.233.116.256	OK	17	2021-02-01T20:31:47+08:00	2021	02	01	12
3	48	221.233.116.256	OK	53	2021-02-01T20:31:48+08:00	2021	02	01	12
4	34	221.233.116.256	OK	32	2021-02-01T20:31:49+08:00	2021	02	01	12
5	50	221.233.116.256	OK	96	2021-02-01T20:31:50+08:00	2021	02	01	12
6	26	221.233.116.256	OK	103	2021-02-01T20:31:53+08:00	2021	02	01	12
7	15	221.233.116.256	OK	108	2021-02-01T20:31:59+08:00	2021	02	01	12
8	35	221.233.116.256	OK	56	2021-02-01T20:32:00+08:00	2021	02	01	12
9	32	221.233.116.256	FAIL	115	2021-02-01T20:32:01+08:00	2021	02	01	12